

Documentation of Condor at USGS CIDA-EMU

By Vladimir Brik and Michael N. Fienen

Open-File Report 2012–xxxx

U.S. Department of the Interior
U.S. Geological Survey

U.S. Department of the Interior
KEN SALAZAR, Secretary

U.S. Geological Survey
Marcia K. McNutt, Director

U.S. Geological Survey, Reston, Virginia 2012

For product and ordering information:
World Wide Web: <http://www.usgs.gov/pubprod>
Telephone: 1-888-ASK-USGS

For more information on the USGS—the Federal source for science about the Earth, its natural and living resources, natural hazards, and the environment:
World Wide Web: <http://www.usgs.gov>
Telephone: 1-888-ASK-USGS

Suggested citation:

Brik, V. and Fienen, M. N., 2012, Documentation of Condor at USGS CIDA-EMU: U.S. Geological Survey Open-File Report 2012–xxx, xx p. [<http://pubs.usgs.gov/of/2011/xxx>]

Any use of trade, product, or firm names in this publication is for descriptive purposes only and does not imply endorsement by the U.S. Government.

Although this report is in the public domain, permission must be secured from the individual copyright owners to reproduce any copyrighted materials contained within this report.

Contents

Figures

Blank page

Documentation of Condor at USGS CIDA-EMU

By Vladimir Brik and Michael N. Fienen

Overview

This document illustrates a Condor framework for running at the USGS CIDA-EMU (Center for Integrated Data Analytics-Environmental Modeling Unit). Two example applications are presented; one using beoPEST (?) and another implementing a leave-one-out-cross validation study (?) .

The first example, beoPEST, is a master-worker code performing model-independent parameter estimation using the techniques of PEST (?). beoPEST is used to distribute the embarrassingly parallel operation of calculating a Jacobian matrix of observation to parameter sensitivities. This operation requires multiple independent model runs, coordinated by a master. The master opens a TCP-IP port on which to listen for available workers. beoPEST is fault-tolerant so workers are allowed to come online and go offline. This configuration is amenable to implementation on Condor. The role of Condor is to start a persistent master and form a queue of potential workers. The master must be exempt from eviction, but workers may be evicted, in which case Condor will start another worker once resources become available. The exemption from eviction for the master can be accomplished either by running on the same machine from which the entire Condor job is submitted, or by specifying a particular target machine on which eviction will not take place. Examples of both processes are presented in this document.

The second example, using leave-one-out cross validation, shows a more general approach to using Condor for a problem in which a problem must be run a large number of times in parallel. A similar approach could be easily applied to a Monte Carlo analysis. In this case, Python code is used to run particular sets of parameters.

Files for Windows beoPEST Run with Local Master

In this example, Python is required to be deployed in addition to the other typically needed files. If Python is not required (or is already deployed on the worker nodes) then "Python27.zip" can be removed from the line of `mw.sub` that starts with `transfer_input_files=` and the first two lines of The executable in this case is a windows batch file called `worker.bat`.

To run this problem, the user must configure the two files `mw.sub` and `worker.bat` (both described below). Then, on the command line, type "`condor_submit mw.sub`". Log files including Condor logs, and both standard out and standard error from beoPEST are available in the `condor_output` subdirectory. The command `condor_q` highlights the activity of the current user while `condor_status` gives a summary of the overall network activity (including which resources are available and which are in use). Finally, to stop a job after it is submitted, type "`condor_rm -all`". This removes all jobs by a current user. Alternatively, single jobs can be killed. The job is referred to as the cluster (and can be identified from `condor_q` output). So, to kill cluster 100, type "`condor_rm 100`".

Condor Submit File to Start Master and Workers: mw.sub

In the following, some preparation is required.

1. In the master folder, a subfolder called `condor_output` must be created.
2. All files needed by beoPEST in slave mode, including PST files, all executables, and all model files, must be zipped into a file called `WORKER_DATA.zip` that resides in the master folder. This folder should initially have been named “data” such that, upon unzipping, a folder called `data` is created. This can be changed if accompanied by changes to `worker.bat`.

In the listings below, all text surrounded by angled braces (<·>) indicates variables that should be replaced with their values in the file. The following are defined:

<casename> The root of the PEST run, such that the control file is <casename>.pst.

<master port> The port on the master machine through which beoPEST master will communicate with slaves.

<requested number of nodes> This is the number of nodes requested. Note that if not enough machines meeting the requirements specified elsewhere in this submit file are available, fewer machines will be used.

```

1 #
2 #Condor Windows submit file for beoPEST with Master and Workers
3 #
4 notification = Never
5 #####
6 # Start the master on the local machine
7 #####
8 universe = local
9 log = condor_output/<casename>_$(Cluster).log
10 output = condor_output/<casename>_$(Cluster).out
11 error = condor_output/<casename>_$(Cluster).err
12 executable = beopest64.exe
13 arguments = <casename> /h :<master port>
14 queue
15 #####
16 # Start the workers
17 #####
18 universe = vanilla
19 PoolName=="CIDA"
20 log = condor_output/w_$(Cluster).log
21 output = condor_output/w_$(Cluster)_$(Process).out
22 error = condor_output/w_$(Cluster)_$(Process).err
23 executable = worker.bat
24 arguments = <casename> $ENV(HOSTNAME) <master port>
25 requirements = ((Target.OpSys=="WINNT61") && (Target.Arch=="INTEL"))
26                && (PoolName=="CIDA") )
27 should_transfer_files = YES

```

4 Documentation of Condor at USGS CIDA-EMU

```
28 when_to_transfer_output = ON_EXIT_OR_EVICT
29 transfer_input_files = unzip.exe,Python27.zip,WORKER_DATA.zip
30 queue <requested number of nodes>
```


Executable to Start Worker: worker.bat

```
1 unzip Python27.zip
2 set path=%cd%\Python27;%cd%;%cd\data;%path%
3
4 unzip WORKER.DATA.zip
5 cd data
6
7 beopest64.exe %1.pst /h %2:%3
```

Files for Linux beoPEST Run

This section presents the raw files being used to run beoPEST in the Condor environment.

Makefile to Prepare Upload and Submit Condor Job: Makefile

A Makefile is used to take advantage of the ability of Make to check status of compressed directories. When Make is called with the argument `submit`, it checks the status of the data and bin folders relative to their compressed versions; if either folder has been changed since the compression was last performed, the compression is performed again and then the two condor jobs are submitted.

```

1 upload: upload/data.tar.gz upload/bin.tar.gz
2
3 upload/data.tar.gz: data/*
4     tar czf upload/data.tar.gz data
5
6 upload/bin.tar.gz: bin/*
7     tar czf upload/bin.tar.gz bin
8
9 submit: upload
10     -mv log/[^_]* log/_archive
11     condor_submit start_master.sub
12     condor_submit start_worker.sub
13     condor_q

```

Condor Submit File to Start Master start_master.sub

```

1 universe = vanilla
2 log = log/master_$(Cluster).log
3 output = log/master_$(Cluster)_$(Process).out
4 error = log/master_$(Cluster)_$(Process).err
5 should_transfer_files = YES
6 requirements = ((Target.Memory>=7000) && (Target.Machine == "a1.hpc"))
7               && (PoolName=="CIDA"))
8 request_memory = 7000
9 executable = job.sh
10 arguments = master <master machine name> <master port> <casename>
11 stream_output = True
12 stream_error = True
13 transfer_output_files = data/<outputfile1>
14 transfer_output_remaps = "<outputfile1> = results/<outputfile1>"
15 when_to_transfer_output = ON_EXIT
16 transfer_input_files = upload/data.tar.gz, upload/bin.tar.gz
17 queue

```

Submit File to Start Worker start_worker.sub

```

1 notification = Never
2
3 universe = vanilla
4 log = log/worker_$(Cluster).log
5 output = log/worker_$(Cluster)_$(Process).out
6 error = log/worker_$(Cluster)_$(Process).err
7 executable = job.sh
8 arguments = worker <master machine name> <master port> <casename>
9 requirements = ((Target.Memory >= 7000) && (PoolName=="CIDA"))
10 request_memory = 7000
11 should_transfer_files = YES
12 when_to_transfer_output = ON_EXIT
13 transfer_input_files = upload/data.tar.gz, upload/bin.tar.gz
14 queue <40>

```

Executable to Start Master or Worker: job.sh

```

1 #!/bin/sh
2
3 role=$1
4 host=$2
5 port=$3
6 casename=$4
7
8 export LD_LIBRARY_PATH=$(pwd)/bin/lib:
9 export PATH=$PATH:$(pwd)/bin
10
11 tar xzf data.tar.gz
12 tar xzf bin.tar.gz
13 #rm data.tar.gz bin.tar.gz
14 cd data
15
16 if [ $role == "worker" ]; then
17     ppest $casename.pst /h $host:$port
18 elif [ $role == "master" ]; then
19     ppest $casename.pst /h :$port
20 else
21     echo "Invalid role <$role>" > /dev/stderr
22 fi

```

Files for Leave One Out Cross Validation

The following files were used to validate a statistical model (?) using Condor and Python.

Makefile to Prepare Upload and Submit Condor Job: Makefile

```

1 upload: upload/LOO.DATA.tgz
2

```

8 Documentation of Condor at USGS CIDA-EMU

```
3 upload/LOO_DATA.tgz: LOO_DATA/*
4     tar czf upload/LOO_DATA.tgz LOO_DATA
5
6 submit: upload
7     -mkdir -p log/_archive
8     -mv log/[^_]* log/_archive
9     condor_submit LOO.sub
10    condor_q
```

Condor Submit File for Leave one out Cross Validation: LOO.sub

```
1 notification = Never
2 universe = vanilla
3 log = log/loo_$(Cluster).log
4 output = log/loo_$(Cluster)_$(Process).out
5 error = log/loo_$(Cluster)_$(Process).err
6 should_transfer_files = YES
7 requirements = ( (OpSys == "LINUX") && (Arch == "X86_64"))
8 executable = worker_loo.sh
9 request_cpus = 1
10 arguments = $(Process)
11 should_transfer_files = yes
12 when_to_transfer_output = ON_EXIT
13 transfer_input_files = upload/LOO_DATA.tgz
14 transfer_output_files = LOO_DATA/data_$(Process).dat
15 queue 99290
```

Executable to Start Leave-one-out Process on Each Condor Node: worker_loo.sh

```
1 #!/bin/sh
2
3 currLooInd=$1
4
5 tar xzf LOO_DATA.tgz
6 rm LOO_DATA.tgz
7
8 export LD_LIBRARY_PATH=.
9 cd LOO_DATA
10
11 ./LOO_single.py $currLooInd
```

Condor Submit File for Leave one out Cross Validation: LOO_single.py

```
1 #!/usr/bin/python
2 import numpy as np
3 import sys
4 from LOO import LOO_Hg
5
6 #####
```

```

7 # Main Function
8 #####
9
10 Masterfile = 'NatfishFinalAllobs_20110617_MNF.csv'
11 Connectionsfile = 'comparable_calcs_20110630.dat'
12 IDS_file = 'all_IDS.dat'
13 allIDS = np.genfromtxt(IDS_file, dtype=int)
14 ind_to_drop = int(sys.argv[1])
15 ID_to_drop = allIDS[ind_to_drop]
16 del allIDS
17
18 cHg, obsHg = LOO.Hg(ID_to_drop, Masterfile, Connectionsfile)
19
20 set1 = open('summaryRESULTS.dat', 'r').readlines()
21
22 ofp = open('data_{0:d}.dat'.format(ind_to_drop), 'w')
23 ofp.write('Dropped_ID—> {0}\n'.format(ID_to_drop))
24 for line in set1:
25     ofp.write(line)
26 ofp.write('%20s%20s%20s\n' % ('dropped_ID', 'modHg', 'measuredHg'))
27 ofp.write('%20d%20.8f%20.8f\n' % (ID_to_drop, cHg[0], obsHg[0]))
28 ofp.close()

```

References Cited

- Brigham, M. E., Fienen, M. N., Donato, D. I., Wente, S. P., Lorenz, D. L., Trombley, M. M., Sanocki, C. A., 2011, Application and validation of the National Descriptive Model of Mercury in Fish (NDMMF), 10th International Conference on Mercury as a Global Pollutant, Halifax, Nova Scotia, Canada, July 25-29, 2011.
- Doherty, John, 2010, PEST–Model-independent parameter estimation–User manual (5th ed., with slight additions): Brisbane, Australia, Watermark Numerical Computing, 336 p.
- Schreüder, W. A., 2009, Running BeoPEST. *in* Proceedings of the 1st PEST Conference, Potomac, Md., 1–3 November.